



# Imitation Learning with F1Tenth

---

Xiatao Sun




# **A Benchmark Comparison of Imitation Learning-based Control Policies for Autonomous Racing**

---

Xiatao Sun, Mingyan Zhou, Zhijun Zhuang, Shuo Yang, Johannes Betz, Rahul Mangharam

# Motivation

- Autonomous racing gains increasing popularity recently
- Most prior works focus on human-engineered or reinforcement learning methods
- Our work applies and compares various imitation learning methods in autonomous racing scenarios




# Demo



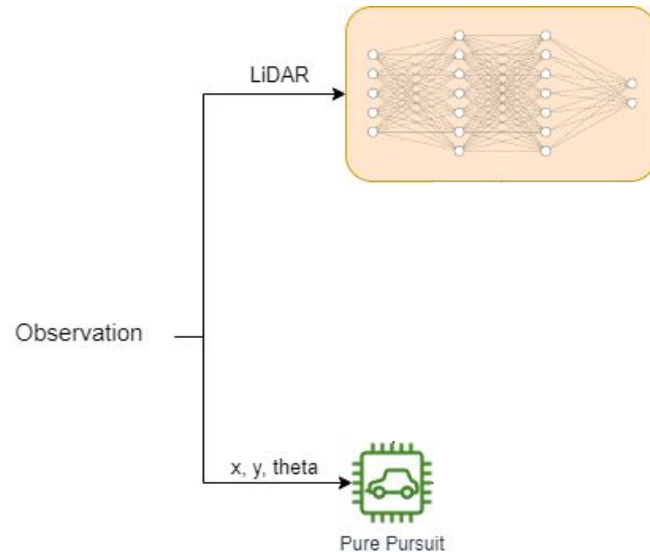
# Implementation

- Expert: pure pursuit algorithm




# Implementation

- Expert: pure pursuit algorithm
- Learner: 2-layer MLP with 256 hidden units



# Implementation

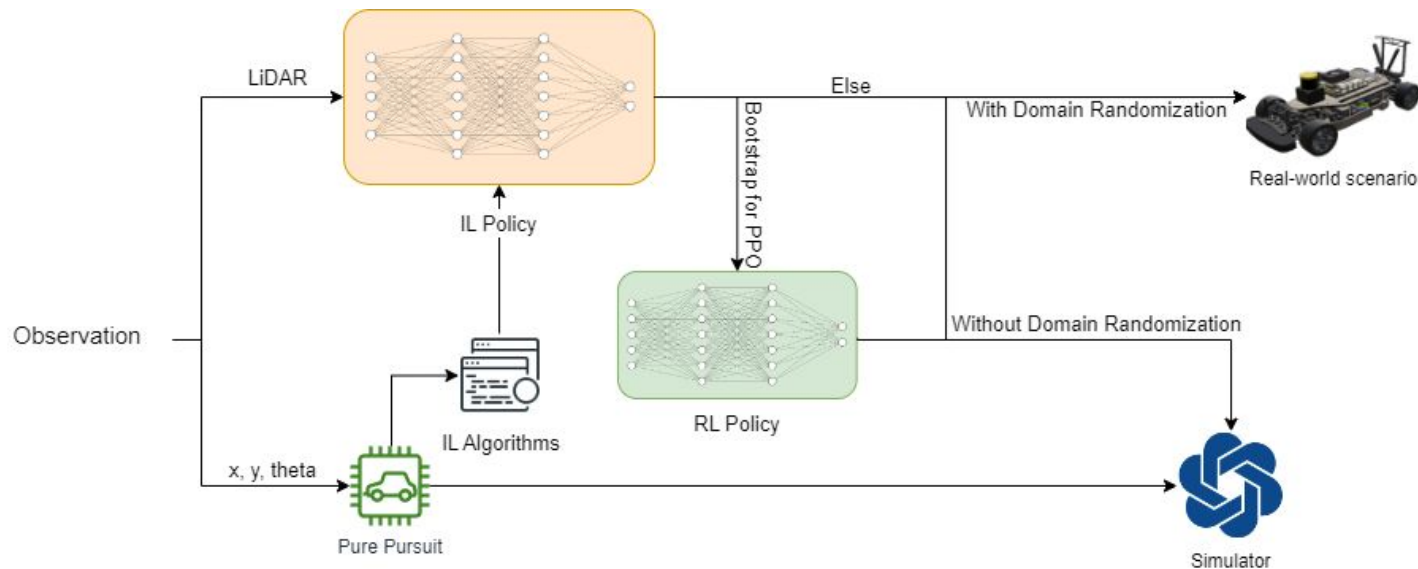
- Expert: pure pursuit algorithm
- Learner: 2-layer MLP with 256 hidden units
- Learning algorithms: BC, DAGGER, HG-DAGGER, EIL



Algorithm	Intervention Rule	Loss Function
EIL	Intervene if $(s, a) \notin \mathcal{G}$	$\ell_B(\cdot) + \lambda \ell_C(\cdot)$
BC	Expert in control	$\ell_C(\cdot)$
DAGGER	Learner in control	$\ell_C(\cdot)^\dagger$
HG-DAGGER	Intervene if $(s, a) \notin \mathcal{G}$	$\ell_C(\cdot)$

# Implementation

- Expert: pure pursuit algorithm
- Learner: 2-layer MLP with 256 hidden units
- Learning algorithms: BC, DAGGER, HG-DAGGER, EIL






# Comparison and Evaluation

---

- Sample and bootstrapping efficiency
- Performance and generalizability
- The experiments are performed in both simulation and real-world environments


# Sample Efficiency

- Training in simulation with respect to 3 different expert
  - Slow: 4.79m/s
  - Normal: 6.39 m/s
  - Fast: 8.24 m/s
- Overall interactive IL demonstrate better sample efficiency



# Bootstrapping efficiency

- Using any IL algorithms can help PPO converge to a better policy
- Interactive IL has better bootstrapping efficiency



# Performance in training map

- Data aggregation mitigates the problem of compounding errors
- EIL has the best performance and is the only one successfully completed one lap when the speed is fast

Expert Type	Expert	BC	DAGGER	HG-DAGGER	EIL
Slow	33.07 s	Failed	34.34 s	33.78 s	33.50 s
Normal	25.04 s	25.35 s	25.85 s	25.06 s	25.22 s
Fast	19.69 s	Failed	Failed	Failed	20.40 s

*Elapsed time of IL policies trained with different experts*

# Generalizability

- The combinations of IL and PPO significantly outperform policies only using IL or PPO, and can efficiently converge to a more generalized policy
- Interactive IL can train policies that are more similar to expert behavior than non-interactive IL and PPO

Method	BC	DAGGER	HG-DAGGER	EIL	PPO	BC+PPO	DAGGER+PPO	HG-DAGGER+PPO	EIL+PPO
Distance Traveled (m)	7.84	8.90	12.34	15.89	12.69	151.23	86.49	155.88	150.15
Complete 1 Lap	No	No	No	No	No	Yes	No	Yes	Yes
Bhattacharyya Distance	0.77	0.60	0.12	0.24	1.09	0.59	0.59	0.47	0.43


*Evaluations of different learned policies in an unseen simulation environment*

# Experiments in real-world environment

- Policies from interactive IL can travel further distances
- PPO policies have more warbling compared with IL policies
- HG-DAGGER + PPO has the best performance

Method	Expert	BC	DAGGER	HG-DAGGER	EIL	PPO	BC+PPO	DAGGER+PPO	HG-DAGGER+PPO	EIL+PPO
Distance Traveled (m)	61.44	6.44	8.49	37.74	38.04	5.27	6.44	6.29	64.08	39.5
Complete 1 Lap	Yes	No	No	No	No	No	No	No	Yes	No

*Evaluations of different learned policies in the real-world environment.*



# Conclusion

---

- IL algorithms can train or bootstrap high-performance policies for autonomous racing
- Interactive mechanism improves sample efficiency and bootstrapping efficiency of IL
- Combining IL and RL can achieve fast convergence and better generalizability



# MEGA-Dagger: Imitation Learning with Multiple Imperfect Experts

---

Xiatao Sun\*, Shuo Yang\*, Rahul Mangharam

\*Authors contributed equally.

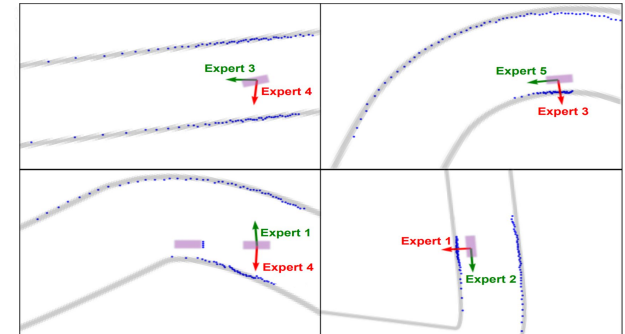


# Overview

- Existing interactive imitation learning algorithms requires **one optimal expert**
- Real-world scenarios: **imperfect experts**
- Challenges:**
  - Undesired demonstrations
  - Conflicted demonstrations



Reduced-scale vehicle controlled by policy learned using HG-Dagger




Example of conflicted demonstrations from different experts

Xiatao Sun, Mingyan Zhou, Zhijun Zhuang, Shuo Yang, Johannes Betz, and Rahul Mangharam. A benchmark comparison of imitation learning-based control policies for autonomous racing. arXiv preprint arXiv:2209.15073, 2022.


Michael Kelly, Chelsea Sidrane, Katherine Driggs-Campbell, and Mykel J Kochenderfer. Hg-dagger: Interactive imitation learning with human experts. In 2019 International Conference on Robotics and Automation (ICRA), pages 8077–8083. IEEE, 2019.

# Comparison: HG-DAgger & MEGA-DAgger

HG-DAgger



MEGA-DAgger (ours)



# Comparison: Experts & MEGA-DAgger

Experts 5



MEGA-DAgger (ours)



# MEGA-Dagger: Multiple Experts GAted DAgger

$\pi_{exp}^j$	Expert policies
$D$	Global dataset
$\pi_{N_0}$	Initial novice policy
$\Pi$	All possible policies
$o$	Observation
$a$	Action
$D_j$	Dataset of current rollout
$\sigma_t$	Safety score
$\pi_{N_i}$	Novice policy of current iteration
$K$	Total number of iterations
$M$	Total number of rollouts for each iteration
$T$	Total number of timesteps

---

## Algorithm 1 MEGA-Dagger

---


```

1: procedure MEGA-DAGGER( $\pi_{exp}^{1:M}$ )
2:   Initialize  $\mathcal{D} \leftarrow \emptyset$ 
3:   Initialize  $\pi_{N_0}$  to any policy in  $\Pi$ 
4:   for iteration  $i = 1 : K$  do
5:     for rollout  $j = 1 : M$  with expert  $\pi_{exp}^j$  do
6:       for timestep  $t \in T$  of rollout  $j$  do
7:         if  $\pi_{exp}^j$  takes control then
8:            $o \leftarrow \text{rollout}_{i,j}^t$ 
9:            $a \leftarrow \pi_{exp}^j(o)$ 
10:           $D_j \leftarrow o, a$ 
11:           $D_j, \sigma_t \leftarrow \text{DATA FILTER}(D_j)$ 
12:        end if
13:      end for
14:       $D_j \leftarrow \text{CONFLICT RESOLUTION}(D_j, D, \sigma_t)$ 
15:       $D \leftarrow D \cup D_j$ 
16:    end for
17:    Train  $\pi_{N_i}$  on  $\mathcal{D}$ 
18:  end for
19: end procedure

```

---

# MEGA-Dagger: Data Filter



---

**Algorithm 1** MEGA-Dagger

---

```
1: procedure MEGA-DAGGER( $\pi_{exp}^{1:M}$ )
2:   Initialize  $\mathcal{D} \leftarrow \emptyset$ 
3:   Initialize  $\pi_{N_0}$  to any policy in  $\Pi$ 
4:   for iteration  $i = 1 : K$  do
5:     for rollout  $j = 1 : M$  with expert  $\pi_{exp}^j$  do
6:       for timestep  $t \in T$  of rollout  $j$  do
7:         if  $\pi_{exp}^j$  takes control then
8:            $o \leftarrow \text{rollout}_{i,j}^t$ 
9:            $a \leftarrow \pi_{exp}^j(o)$ 
10:           $D_j \leftarrow o, a$ 
11:           $D_j, \sigma_t \leftarrow \text{DATA FILTER}(D_j)$ 
12:        end if
13:      end for
14:       $D_j \leftarrow \text{CONFLICT RESOLUTION}(D_j, D, \sigma_t)$ 
15:       $D \leftarrow D \cup D_j$ 
16:    end for
17:    Train  $\pi_{N_i}$  on  $\mathcal{D}$ 
18:  end for
19: end procedure
```

---

# MEGA-Dagger: Data Filter

- Control Barrier Function (CBF):

$$h(x_t^e, y_t^e) = (x_t^e - x_t^p)^2 - (y_t^e - y_t^p)^2 - \alpha^2$$

- Safety score:

$$\sigma_t = h(x_{t+1}^e, y_{t+1}^e) - (1 - \gamma)h(x_t^e, y_t^e)$$

$$0 < \gamma \leq 1$$

$(x_t^e, y_t^e)$  Current ego position

$(x_t^p, y_t^p)$  Current obstacle position

$\alpha$  Minimal safety distance

---

**Algorithm 1** MEGA-Dagger


---

```
1: procedure MEGA-DAGGER( $\pi_{exp}^{1:M}$ )
2:   Initialize  $\mathcal{D} \leftarrow \emptyset$ 
3:   Initialize  $\pi_{N_0}$  to any policy in  $\Pi$ 
4:   for iteration  $i = 1 : K$  do
5:     for rollout  $j = 1 : M$  with expert  $\pi_{exp}^j$  do
6:       for timestep  $t \in T$  of rollout  $j$  do
7:         if  $\pi_{exp}^j$  takes control then
8:            $o \leftarrow \text{rollout}_{i,j}^t$ 
9:            $a \leftarrow \pi_{exp}^j(o)$ 
10:           $D_j \leftarrow o, a$ 
11:           $D_j, \sigma_t \leftarrow \text{DATA FILTER}(D_j)$ 
12:        end if
13:      end for
14:       $D_j \leftarrow \text{CONFLICT RESOLUTION}(D_j, D, \sigma_t)$ 
15:       $D \leftarrow D \cup D_j$ 
16:    end for
17:    Train  $\pi_{N_i}$  on  $\mathcal{D}$ 
18:  end for
19: end procedure
```

---

# Evaluation: Data Filter


- Comparison: HG-Dagger with data filter and HG-Dagger with random data truncation
- The proposed data filter is able to effectively remove undesired demonstrations and result in a better policy



- $P(U)$ : Undesired behavior probability
- $r_\beta$ : Ratio of removed demonstrations

$P(U)$	0.1	0.2	0.3	0.4	0.5
$r_\beta$	0.22	0.25	0.28	0.32	0.34
$P(U)$	0.6	0.7	0.8	0.9	1.0
$r_\beta$	0.38	0.49	0.57	0.69	0.82

# MEGA-Dagger: Conflict Resolution



---

**Algorithm 1** MEGA-Dagger

---

```
1: procedure MEGA-DAGGER( $\pi_{exp}^{1:M}$ )
2:   Initialize  $\mathcal{D} \leftarrow \emptyset$ 
3:   Initialize  $\pi_{N_0}$  to any policy in  $\Pi$ 
4:   for iteration  $i = 1 : K$  do
5:     for rollout  $j = 1 : M$  with expert  $\pi_{exp}^j$  do
6:       for timestep  $t \in T$  of rollout  $j$  do
7:         if  $\pi_{exp}^j$  takes control then
8:            $o \leftarrow \text{rollout}_{i,j}^t$ 
9:            $a \leftarrow \pi_{exp}^j(o)$ 
10:           $D_j \leftarrow o, a$ 
11:           $D_j, \sigma_t \leftarrow \text{DATA FILTER}(D_j)$ 
12:        end if
13:      end for
14:       $D_j \leftarrow \text{CONFLICT RESOLUTION}(D_j, D, \sigma_t)$ 
15:       $D \leftarrow D \cup D_j$ 
16:    end for
17:    Train  $\pi_{N_i}$  on  $\mathcal{D}$ 
18:  end for
19: end procedure
```

---



# MEGA-Dagger: Conflict Resolution

- Cosine similarities:

$$\Theta = \frac{O \cdot O_j}{\|O\| \odot \|O_j\|}$$

- Evaluation score:

$$\omega_t = \frac{\|\sigma_t\| - \min_t \|\sigma_t\|}{\max_t \|\sigma_t\| - \min_t \|\sigma_t\|} + \frac{\|v_t\| - \min_t \|v_t\|}{\max_t \|v_t\| - \min_t \|v_t\|}$$

$O$  Observations in  $D$

$O_j$  Observations in  $D_j$

$v_t$  Ego velocity

---

**Algorithm 1** MEGA-Dagger

---

```
1: procedure MEGA-DAGGER( $\pi_{exp}^{1:M}$ )
2:   Initialize  $\mathcal{D} \leftarrow \emptyset$ 
3:   Initialize  $\pi_{N_0}$  to any policy in  $\Pi$ 
4:   for iteration  $i = 1 : K$  do
5:     for rollout  $j = 1 : M$  with expert  $\pi_{exp}^j$  do
6:       for timestep  $t \in T$  of rollout  $j$  do
7:         if  $\pi_{exp}^j$  takes control then
8:            $o \leftarrow \text{rollout}_{i,j}^t$ 
9:            $a \leftarrow \pi_{exp}^j(o)$ 
10:           $D_j \leftarrow o, a$ 
11:           $D_j, \sigma_t \leftarrow \text{DATA FILTER}(D_j)$ 
12:        end if
13:      end for
14:       $D_j \leftarrow \text{CONFLICT RESOLUTION}(D_j, D, \sigma_t)$ 
15:       $D \leftarrow D \cup D_j$ 
16:    end for
17:    Train  $\pi_{N_i}$  on  $\mathcal{D}$ 
18:  end for
19: end procedure
```

---

# Evaluation: Comparison with HG-DAgger

- Comparison: MEGA-DAgger, HG-DAgger with and without data filter
- Perform experiments on two different maps
- Policies learned using MEGA-DAgger demonstrates better performance on both overtaking and collision avoidance!

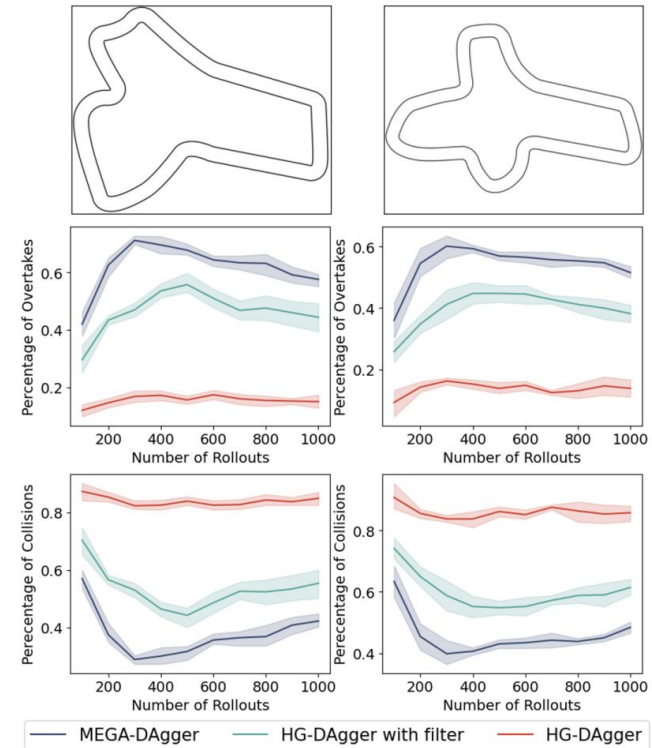
**Overtakes rate**

**Collisions rate**

— MEGA-DAgger  
— HG-DAgger with filter  
— HG-DAgger

**Map 1**

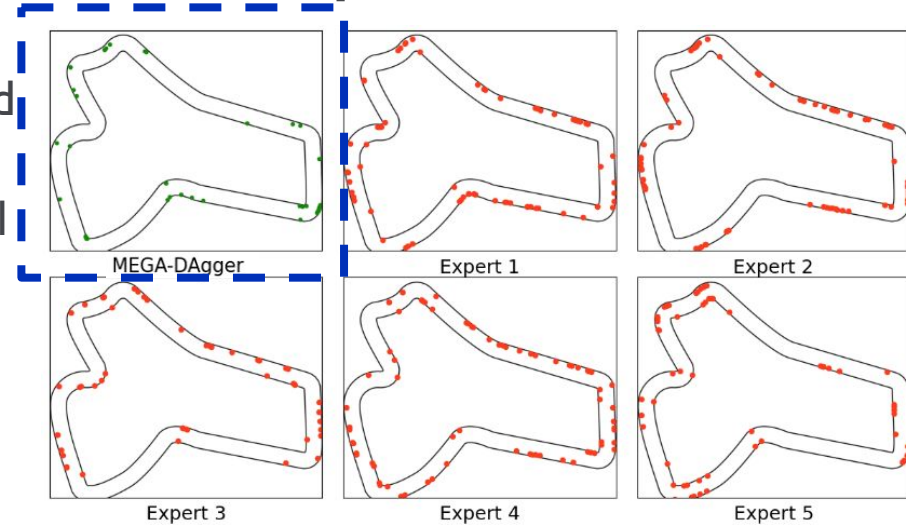
**Map 2**



# Evaluation: Comparison with Experts

- Comparison: MEGA-Dagger and experts
- By leveraging the complementarily good demonstrations, policies learned using MEGA-Dagger is able to outperform all experts!


## Collision points visualization



Metrics	MEGA-Dagger	Expert 1	Expert 2	Expert 3	Expert 4	Expert 5	Experts Cumulative
Collisions Percentage	<b><math>0.212 \pm 0.019</math></b>	$0.340 \pm 0.025$	$0.401 \pm 0.033$	$0.291 \pm 0.025$	$0.392 \pm 0.028$	$0.317 \pm 0.032$	$0.348 \pm 0.051$
Overtakes Percentage	<b><math>0.781 \pm 0.016</math></b>	$0.657 \pm 0.027$	$0.594 \pm 0.036$	$0.706 \pm 0.022$	$0.605 \pm 0.024$	$0.681 \pm 0.030$	$0.649 \pm 0.051$



# Thank you



Source code: <https://github.com/M4D-SCIENTIST/MEGA-Dagger>